

Stock Market Prediction Using Deep Reinforcement learning:

Ensemble strategy

2.1. Install all the packages through FinRL library

```
## install finrl library
!pip install FinRL

from finrl import config
from finrl import config_tickers
import os
if not os.path.exists("./" + config.DATA_SAVE_DIR):
    os.makedirs("./" + config.DATA_SAVE_DIR)
if not os.path.exists("./" + config.TRAINED_MODEL_DIR):
    os.makedirs("./" + config.TRAINED_MODEL_DIR)
if not os.path.exists("./" + config.TENSORBOARD_LOG_DIR):
    os.makedirs("./" + config.TENSORBOARD_LOG_DIR)
if not os.path.exists("./" + config.RESULTS_DIR):
    os.makedirs("./" + config.RESULTS_DIR)
```

2.2. Import Packages

```
import pandas as pd
import numpy as np
import matplotlib
import matplotlib.pyplot as plt
# matplotlib.use('Agg')
import datetime

%matplotlib inline
from finrl.finrl_meta.preprocessor.yahoodownloader import YahooDownloader
from finrl.finrl_meta.preprocessor.preprocessors import FeatureEngineer, data_split
from finrl.finrl_meta.env_stock_trading.env_stocktrading import StockTradingEnv
from finrl.agents.stablebaselines3.models import DRLAgent
from finrl.finrl_meta.data_processor import DataProcessor

from finrl.plot import backtest_stats, backtest_plot, get_daily_return, get_baseline
from pprint import pprint

import sys
sys.path.append("../FinRL-Library")
```

```
import itertools
```

```
/usr/local/lib/python3.7/dist-packages/pyfolio/pos.py:27: UserWarning: Module "zipline.assets" not found; multipliers will not be applied'
```

▼ Part 3. Download Data

```
# from config.py TRAIN_START_DATE is a string
config.TRAIN_START_DATE
```

```
# from config.py TRAIN_END_DATE is a string
config.TRAIN_END_DATE
```

```
df = YahooDownloader(start_date = '2009-01-01',
                     end_date = '2021-10-31',
                     ticker_list = config_tickers.DOW_30_TICKER).fetch_data()
```

[illegible]

```
Shape of DataFrame: (94331, 8)
```

```
print(config_tickers.DOW_30_TICKER)
```

```
['AXP', 'AMGN', 'AAPL', 'BA', 'CAT', 'CSCO', 'CVX', 'GS', 'HD', 'HON', 'IBM', 'INTC', 'JNJ', 'KO', 'LLY', 'MRK', 'MSFT', 'NKE', 'PFE', 'PG', 'PYPL', 'QCOM', 'ROST', 'T', 'TSLA', 'UNH', 'V', 'VZ', 'WMT', 'XOM']
```

```
df.shape
```

```
(94331, 8)
```

```
df.sort_values(['date', 'tic'], ignore_index=True).head()
```

	date	open	high	low	close	volume	tic	day
0	2009-01-02	3.067143	3.251429	3.041429	2.771174	746015200	AAPL	4
1	2009-01-02	58.590000	59.080002	57.750000	44.867592	6547900	AMGN	4
2	2009-01-02	18.570000	19.520000	18.400000	15.535348	10955700	AXP	4
3	2009-01-02	42.799999	45.560001	42.779999	33.941101	7010200	BA	4
4	2009-01-02	44.910000	46.980000	44.709999	32.164726	7117200	CAT	4

▼ Part 4: Preprocess Data

```
fe = FeatureEngineer(
    use_technical_indicator=True,
    tech_indicator_list = config.INDICATORS,
    use_vix=True,
    use_turbulence=True,
    user_defined_feature = False)
```

```
processed = fe.preprocess_data(df)
```

```
Successfully added technical indicators
[*****100%*****] 1 of 1 completed
Shape of DataFrame: (3229, 8)
Successfully added vix
Successfully added turbulence index
Successfully added technical indicators
[*****100%*****] 1 of 1 completed
Shape of DataFrame: (3229, 8)
Successfully added vix
Successfully added turbulence index
```

```
list_ticker = processed["tic"].unique().tolist()
list_date = list(pd.date_range(processed['date'].min(), processed['date'].max()).astype(str))
```

```

combination = list(itertools.product(list_date,list_ticker))

processed_full = pd.DataFrame(combination,columns=["date","tic"]).merge(processed,on=["date",
processed_full = processed_full[processed_full['date'].isin(processed['date'])]
processed_full = processed_full.sort_values(['date','tic'])

processed_full = processed_full.fillna(0)

processed_full.sort_values(['date','tic'],ignore_index=True).head(10)

```

	date	tic	open	high	low	close	volume	day	macd	boll
0	2009-01-02	AAPL	3.067143	3.251429	3.041429	2.771174	746015200.0	4.0	0.0	2.99
1	2009-01-02	AMGN	58.590000	59.080002	57.750000	44.867592	6547900.0	4.0	0.0	2.99
2	2009-01-02	AXP	18.570000	19.520000	18.400000	15.535348	10955700.0	4.0	0.0	2.99
3	2009-01-02	BA	42.799999	45.560001	42.779999	33.941101	7010200.0	4.0	0.0	2.99
4	2009-01-02	CAT	44.910000	46.980000	44.709999	32.164726	7117200.0	4.0	0.0	2.99
5	2009-01-02	CRM	8.025000	8.550000	7.912500	8.505000	4069200.0	4.0	0.0	2.99

▼ Part 5. Design Environment

```

train = data_split(processed_full, '2009-01-01','2020-07-01')
trade = data_split(processed_full, '2020-07-01','2021-10-31')
print(len(train))
print(len(trade))

83897
9744

train.tail()

```

	date	tic	open	high	low	close	volume	day	ma
2892	2020-06-30	UNH	288.570007	296.450012	287.660004	287.776794	2932900.0	1.0	-0.0194
2892	2020-06-30	V	191.490005	193.750000	190.160004	190.737213	9040100.0	1.0	1.0487
2892	2020-06-30	VZ	54.919998	55.290001	54.360001	50.376743	17414800.0	1.0	-0.4371
2892	2020-06-30	WBA	42.119999	42.580002	41.759998	39.035736	4782100.0	1.0	-0.0831
2892	2020-06-30	WMT	119.220001	120.129997	118.540001	116.121773	6836400.0	1.0	-0.8861

```
trade.head()
```

	date	tic	open	high	low	close	volume	day	ma
0	2020-07-01	AAPL	91.279999	91.839996	90.977501	89.904602	110737200.0	2.0	3.01460
0	2020-07-01	AMGN	235.520004	256.230011	232.580002	240.153961	6575800.0	2.0	3.63639
0	2020-07-01	AXP	95.250000	96.959999	93.639999	92.086372	3301000.0	2.0	-0.38916
0	2020-07-01	BA	185.880005	190.610001	180.039993	180.320007	49036700.0	2.0	5.44319
0	2020-07-01	CAT	129.380005	129.399994	125.879997	120.651642	2807800.0	2.0	1.27262

```
config.INDICATORS
```

```
[ 'macd',
  'boll_ub',
  'boll_lb',
  'rsi_30',
  'cci_30',
  'dx_30',
  'close_30_sma',
  'close_60_sma' ]
```

```
stock_dimension = len(train.tic.unique())
state_space = 1 + 2*stock_dimension + len(config.INDICATORS)*stock_dimension
```

```
print(f"Stock Dimension: {stock_dimension}, State Space: {state_space}")
```

```
    Stock Dimension: 29, State Space: 291
```

```
buy_cost_list = sell_cost_list = [0.001] * stock_dimension
num_stock_shares = [0] * stock_dimension
```

```
env_kwargs = {
    "hmax": 100,
    "initial_amount": 1000000,
    "num_stock_shares": num_stock_shares,
    "buy_cost_pct": buy_cost_list,
    "sell_cost_pct": sell_cost_list,
    "state_space": state_space,
    "stock_dim": stock_dimension,
    "tech_indicator_list": config.INDICATORS,
    "action_space": stock_dimension,
    "reward_scaling": 1e-4
}
```

```
e_train_gym = StockTradingEnv(df = train, **env_kwargs)
```

▼ Environment for Training

```
env_train, _ = e_train_gym.get_sb_env()
print(type(env_train))
```

```
<class 'stable_baselines3.common.vec_env.dummy_vec_env.DummyVecEnv'>
```

▼ Part 6: Implement DRL Algorithms

```
agent = DRLAgent(env = env_train)
```

Model Training: 3 models, A2C DDPG, PPO

▼ Model 1: A2C

```
agent = DRLAgent(env = env_train)
model_a2c = agent.get_model("a2c")
```

```
{'n_steps': 5, 'ent_coef': 0.01, 'learning_rate': 0.0007}
Using cuda device
```

```
trained_a2c = agent.train_model(model=model_a2c,
                                tb_log_name='a2c',
                                total_timesteps=50000)
```

time/		
fps		97
iterations		100
time_elapsed		5
total_timesteps		500
train/		
entropy_loss		-41.3
explained_variance		0.461
learning_rate		0.0007
n_updates		99
policy_loss		-19.4
reward		0.18487218
std		1
value_loss		0.456

time/		
fps		100
iterations		200
time_elapsed		9
total_timesteps		1000
train/		
entropy_loss		-41.3
explained_variance		-0.0554
learning_rate		0.0007
n_updates		199
policy_loss		-61.1
reward		-0.18095054
std		1.01
value_loss		2.58

time/		
fps		100
iterations		300
time_elapsed		14
total_timesteps		1500
train/		
entropy_loss		-41.3
explained_variance		0
learning_rate		0.0007
n_updates		299
policy_loss		-238
reward		4.376703
std		1.01
value_loss		31.3

time/		
fps	100	
iterations	400	
time_elapsed	19	
total_timesteps	2000	
train/		
entropy_loss	-41.3	
explained_variance	-1.19e-07	
learning_rate	0.0007	

▼ Model 2: DDPG

```
agent = DRLAgent(env = env_train)
model_ddpg = agent.get_model("ddpg")

{'batch_size': 128, 'buffer_size': 50000, 'learning_rate': 0.001}
Using cuda device
```

```
trained_ddpg = agent.train_model(model=model_ddpg,
                                tb_log_name='ddpg',
                                total_timesteps=50000)
```

```
day: 2892, episode: 20
begin_total_asset: 1000000.00
end_total_asset: 4774133.37
total_reward: 3774133.37
total_cost: 5470.92
total_trades: 47732
Sharpe: 0.868
```

```
=====
```

time/		
episodes	4	
fps	71	
time_elapsed	162	
total_timesteps	11572	
train/		
actor_loss	96.6	
critic_loss	1.47e+03	
learning_rate	0.001	
n_updates	8679	
reward	2.8689466	

time/		
episodes	8	
fps	66	
time_elapsed	347	
total_timesteps	23144	
train/		
actor_loss	24.9	
critic_loss	14.5	

learning_rate	0.001
n_updates	20251
reward	2.8689466

```

-----
day: 2892, episode: 30
begin_total_asset: 1000000.00
end_total_asset: 4390924.34
total_reward: 3390924.34
total_cost: 999.00
total_trades: 40488
Sharpe: 0.730
=====

```

time/	
episodes	12
fps	65
time_elapsed	532
total_timesteps	34716
train/	
actor_loss	6.26
critic_loss	3.49
learning_rate	0.001
n_updates	31823
reward	2.8689466

time/	
episodes	16

▼ Model 3: PPO

```

agent = DRLAgent(env = env_train)
PPO_PARAMS = {
    "n_steps": 2048,
    "ent_coef": 0.01,
    "learning_rate": 0.00025,
    "batch_size": 128,
}
model_ppo = agent.get_model("ppo", model_kwargs = PPO_PARAMS)

{'n_steps': 2048, 'ent_coef': 0.01, 'learning_rate': 0.00025, 'batch_size': 128}
Using cuda device

```

```

trained_ppo = agent.train_model(model=model_ppo,
                                tb_log_name='ppo',
                                total_timesteps=50000)

```

time/	
fps	115
iterations	1
time_elapsed	17

total_timesteps	2048
train/	
reward	0.10877045

time/	
fps	112
iterations	2
time_elapsed	36
total_timesteps	4096
train/	
approx_kl	0.015766323
clip_fraction	0.226
clip_range	0.2
entropy_loss	-41.2
explained_variance	-0.0621
learning_rate	0.00025
loss	2.75
n_updates	10
policy_gradient_loss	-0.0268
reward	0.44184494
std	1
value_loss	9.5

day: 2892, episode: 40
 begin_total_asset: 1000000.00
 end_total_asset: 2858099.69
 total_reward: 1858099.69
 total_cost: 335985.43
 total_trades: 80095
 Sharpe: 0.673

time/	
fps	111
iterations	3
time_elapsed	55
total_timesteps	6144
train/	
approx_kl	0.014054896
clip_fraction	0.127
clip_range	0.2
entropy_loss	-41.2
explained_variance	-0.000786
learning_rate	0.00025
loss	30.3
n_updates	20
policy_gradient_loss	-0.0177
reward	-1.5774492
std	1
value_loss	54.6

▼ Trading

Assume that we have \$1.000.000 initial capital at 2020-07-01. We use the DDPG model to trade.

```
data_risk_indicator = processed_full[(processed_full.date<'2020-07-01') & (processed_full.date>'2021-10-31')]
insample_risk_indicator = data_risk_indicator.drop_duplicates(subset=['date'])
```

```
insample_risk_indicator.vix.describe()
```

```
count    2893.000000
mean      18.824245
std       8.489311
min       9.140000
25%      13.330000
50%      16.139999
75%      21.309999
max      82.690002
Name: vix, dtype: float64
```

```
insample_risk_indicator.vix.quantile(0.996)
```

```
57.40400183105453
```

```
insample_risk_indicator.turbulence.describe()
```

```
count    2893.000000
mean      34.567961
std      43.790810
min       0.000000
25%      14.962718
50%      24.124747
75%      39.162552
max     652.506566
Name: turbulence, dtype: float64
```

```
insample_risk_indicator.turbulence.quantile(0.996)
```


```
276.45175727554096
```

▼ Trade

```
#trade = data_split(processed_full, '2020-07-01','2021-10-31')
e_trade_gym = StockTradingEnv(df = trade, turbulence_threshold = 70,risk_indicator_col='vix',
```

```
trade.head()
```

	date	tic	open	high	low	close	volume	day	max
0	2020-07-01	AAPL	91.279999	91.839996	90.977501	89.904602	110737200.0	2.0	3.01460
0	2020-07-01	AMGN	235.520004	256.230011	232.580002	240.153961	6575800.0	2.0	3.63639
0	2020-07-01	AXP	95.250000	96.959999	93.639999	92.086372	3301000.0	2.0	-0.38916
0	2020-07-01	BA	185.880005	190.610001	180.039993	180.320007	49036700.0	2.0	5.44319
0	2020-07-01	CAT	129.380005	129.399994	125.879997	120.651642	2807800.0	2.0	1.27262



```
df_account_value, df_actions = DRLAgent.DRL_prediction(
    model=trained_ddpg,
    environment = e_trade_gym)
```

```
hit end!
```

```
df_account_value.shape
```

```
(336, 2)
```

```
df_account_value.tail()
```

	date	account_value
331	2021-10-22	1.424764e+06
332	2021-10-25	1.433647e+06
333	2021-10-26	1.427543e+06
334	2021-10-27	1.426441e+06
335	2021-10-28	1.440410e+06

```
df_actions.head()
```

	AAPL	AMGN	AXP	BA	CAT	CRM	CSCO	CVX	DIS	GS	...	MRK	MSFT	NKE	PG	TRV	I
date																	
2020-07-01	0	0	0	98	99	0	0	0	0	77	...	0	95	0	0	0	
2022																	

▼ Part 7: Backtest Our Strategy

▼ 7.1 BackTestStats

```
print("====Get Backtest Results====")
now = datetime.datetime.now().strftime('%Y%m%d-%H%M')

perf_stats_all = backtest_stats(account_value=df_account_value)
perf_stats_all = pd.DataFrame(perf_stats_all)
perf_stats_all.to_csv("./"+config.RESULTS_DIR+"/perf_stats_all_"+now+'.csv')
```

```
====Get Backtest Results====
```

```
Annual return      0.314815
Cumulative returns  0.440410
Annual volatility   0.135542
Sharpe ratio        2.094102
Calmar ratio        4.162814
Stability           0.952226
Max drawdown        -0.075625
Omega ratio         1.407491
Sortino ratio        3.063127
Skew                NaN
Kurtosis            NaN
Tail ratio          1.091429
Daily value at risk -0.015950
dtype: float64
```

```
#baseline stats
```

```
print("====Get Baseline Stats====")
baseline_df = get_baseline(
    ticker="^DJI",
    start = df_account_value.loc[0,'date'],
    end = df_account_value.loc[len(df_account_value)-1,'date'])
```

```
stats = backtest_stats(baseline_df, value_col_name = 'close')
```

```
====Get Baseline Stats====
```

```
[*****100%*****] 1 of 1 completed
Shape of DataFrame: (335, 8)
```

Annual return	0.273520
Cumulative returns	0.379084
Annual volatility	0.139248
Sharpe ratio	1.811893
Calmar ratio	3.062662
Stability	0.918651
Max drawdown	-0.089308
Omega ratio	1.351851
Sortino ratio	2.684720
Skew	NaN
Kurtosis	NaN
Tail ratio	1.051856
Daily value at risk	-0.016542

dtype: float64

```
df_account_value.loc[0, 'date']
```

```
'2020-07-01'
```

```
df_account_value.loc[len(df_account_value)-1, 'date']
```

```
'2021-10-28'
```

▼ 7.2 BackTestPlot

```
print("=====Compare to DJIA=====")
%matplotlib inline
# Dow Jones Index: ^DJI
backtest_plot(df_account_value,
               baseline_ticker = '^DJI',
               baseline_start = df_account_value.loc[0, 'date'],
               baseline_end = df_account_value.loc[len(df_account_value)-1, 'date'])
```

=====Compare to DJIA=====

[*****100%*****] 1 of 1 completed

Shape of DataFrame: (335, 8)

Start date 2020-07-01

End date 2021-10-28

Total months 16

Backtest

Annual return 31.481%

Cumulative returns 44.041%

Annual volatility 13.554%

Sharpe ratio 2.09

Calmar ratio 4.16

Stability 0.95

Max drawdown -7.563%

Omega ratio 1.41

Sortino ratio 3.06

Skew NaN

Kurtosis NaN

Tail ratio 1.09

Daily value at risk -1.595%

Alpha 0.07

Beta 0.86

Worst drawdown periods	Net drawdown in %	Peak date	Valley date	Recovery date	Durat
0	7.56	2020-09-02	2020-09-23	2020-10-13	

▼ Part 6: Implement DRL Algorithms

```

rebalance_window = 63
validation_window = 63
train_start = '2000-01-01'
train_end = '2019-01-01'
val_test_start = '2019-01-01'
val_test_end = '2021-01-18'

```

```

ensemble_agent = DRLEnsembleAgent(df=processed_full,
                                   train_period=(train_start,train_end),

```

```

val_test_period=(val_test_start,val_test_end),
rebalance_window=rebalance_window,
validation_window=validation_window,
**env_kwargs)

```

```

A2C_model_kwargs = {
    'n_steps': 5,
    'ent_coef': 0.01,
    'learning_rate': 0.0005
}

```

```

PPO_model_kwargs = {
    "ent_coef":0.01,
    "n_steps": 2048,
    "learning_rate": 0.00025,
    "batch_size": 128
}

```

```

DDPG_model_kwargs = {
    "action_noise":"ornstein_uhlenbeck",
    "buffer_size": 50_000,
    "learning_rate": 0.000005,
    "batch_size": 128
}

```

```

timesteps_dict = {'a2c' : 100_000,
                  'ppo' : 100_000,
                  'ddpg' : 50_000
}

```

```

df_summary = ensemble_agent.run_ensemble_strategy(A2C_model_kwargs,
                                                  PPO_model_kwargs,
                                                  DDPG_model_kwargs,
                                                  timesteps_dict)

```

```

=====Start Ensemble Strategy=====

```

```

=====

```

```

nan

```

```

turbulence_threshold: 5625265.115256409

```

```

=====Model training from: 2000-01-01 to 2019-01-02

```

```

=====A2C Training=====

```

```

{'n_steps': 5, 'ent_coef': 0.01, 'learning_rate': 0.0005}

```

```

Using cpu device

```

```

Logging to tensorboard_log/a2c/a2c_126_1

```

```

-----
| time/          |          |
|   fps          |    99    |
|  iterations    |   100    |
| time_elapsed   |    5     |
| total_timesteps|   500    |
| train/         |          |
| entropy_loss   |  -35.6   |

```


explained_variance	0.229
learning_rate	0.0005
n_updates	99
policy_loss	-77
std	1
value_loss	4.79

time/	
fps	99
iterations	200
time_elapsed	10
total_timesteps	1000
train/	
entropy_loss	-35.6
explained_variance	-1.19e-07
learning_rate	0.0005
n_updates	199
policy_loss	-40.5
std	1.01
value_loss	2.02

time/	
fps	99
iterations	300
time_elapsed	15
total_timesteps	1500
train/	
entropy_loss	-35.6
explained_variance	0.188
learning_rate	0.0005
n_updates	299
policy_loss	250
std	1.01
value_loss	59.4

time/	
fps	99

df_summary

	Iter	Val Start	Val End	Model Used	A2C Sharpe	PPO Sharpe	DDPG Sharpe
0	126	2019-01-02	2019-04-01	DDPG	-0.230842	-0.0926179	0.320516

▼ Part 7: Backtest Our Strategy

```

unique_trade_date = processed_full[(processed_full.date > val_test_start)&(processed_full.date < val_test_end)]

df_trade_date = pd.DataFrame({'date':unique_trade_date})

df_account_value=pd.DataFrame()
for i in range(rebalance_window+validation_window, len(unique_trade_date)+1,rebalance_window):
    temp = pd.read_csv('results/account_value_trade_{}_{}.csv'.format('ensemble',i))
    df_account_value = df_account_value.append(temp,ignore_index=True)
sharpe=(252**0.5)*df_account_value.account_value.pct_change(1).mean()/df_account_value.account_value.std()
print('Sharpe Ratio: ',sharpe)
df_account_value=df_account_value.join(df_trade_date[validation_window:].reset_index(drop=True))

Sharpe Ratio: 0.17377457928215234

df_account_value.head()

```

	account_value	date	daily_return	date
0	500000.000000	2019-04-01	NaN	2019-04-01
1	498557.292308	2019-04-02	-0.002885	2019-04-02
2	497280.859251	2019-04-03	-0.002560	2019-04-03
3	515844.876585	2019-04-04	0.037331	2019-04-04
4	496792.628660	2019-04-05	-0.036934	2019-04-05

▼ 7.1 BackTestPlot

```

%matplotlib inline
BackTestPlot(df_account_value,
              baseline_ticker = '^DJI',
              baseline_start = df_account_value.loc[0,'date'],
              baseline_end = df_account_value.loc[len(df_account_value)-1,'date'])

```

Start date	2016-01-04
End date	2020-05-12
Total months	52
Backtest	
Annual return	13.162%
Cumulative returns	71.305%
Annual volatility	7.868%
Sharpe ratio	1.61
Calmar ratio	2.28
Stability	0.90
Max drawdown	-5.772%
Omega ratio	1.46
Sortino ratio	2.63
Skew	NaN
Kurtosis	NaN
Tail ratio	1.47
Daily value at risk	-0.941%
Alpha	0.11
Beta	0.14

	Worst drawdown periods	Net drawdown in %	Peak date	Valley date	Recovery date	Duration
0		5.77	2016-02-01	2016-02-11	2016-02-25	19
1		5.00	2016-07-14	2016-10-17	2016-11-10	86
2		4.96	2017-03-01	2017-04-13	2017-10-04	156
3		4.34	2018-06-12	2018-06-28	2019-02-04	170
4		3.85	2016-06-23	2016-06-27	2016-06-30	6

/Users/hongyangyang/anaconda3/lib/python3.6/site-packages/pyfolio/tears.py:907: UserWarning: 'interesting times.', UserWarning)

